



IMapBook Collaborative Discussion Classification

Uroš Polanc

Abstract

Automatic text classification has been considered as a vital method to manage and process a vast amount of documents in digital form, which is widespread and continuously increasing. That brought a wave of new tools for student education, such as the IMapBook platform. It gives students a way to communicate and discuss the books they are reading. In our assignment, we will try to classify those discussions based on the book it's relevant to, and the type of message.

Keywords

IMapBook, Text Classification, TF-IDF, Handcrafted Features, BERT, Logistic Regression

Advisors : Slavko Žitnik

Introduction

With the rapid development of the internet and big data, digital information is increasing at a high rate. With this technologies and applications such as IMapBook [1], a web-based application that allows for discussions on reading materials as well as interactive games, also started emerging.

In our assignment, we will look at some IMapBook collaborative discussions on different books. For that, we first need to look at some vital information, that we extracted from the discussions and the book texts. We will create two different message classifications, based on :

- Book ID (3-class),
- Code Preliminary (16-class).

From the message types we can say that not all are important to the discussions of the books, and only the content discussion and content question are directly related to the books.

With this we were also provided with a (1) dataset that contains all message types and (2) dataset that contains only messages important to the discussion. With this we will be able to further compare how much noise impacts the classifiers and prediction.

By checking the manual classifications in the dataset, we noticed some were misspelled and some were incorrectly written. Therefore, we applied some fixes to the manual

classifications. The errors such as excess whitespace, lower and upper case, etc... were removed. Classifications where there were multiple classes, were simplified into one class, to simplify the classification process. The final fixed distribution is not intraclass uniform as can be seen in figures 1 and 2.




Figure 1. The distribution of book IDs in the dataset.

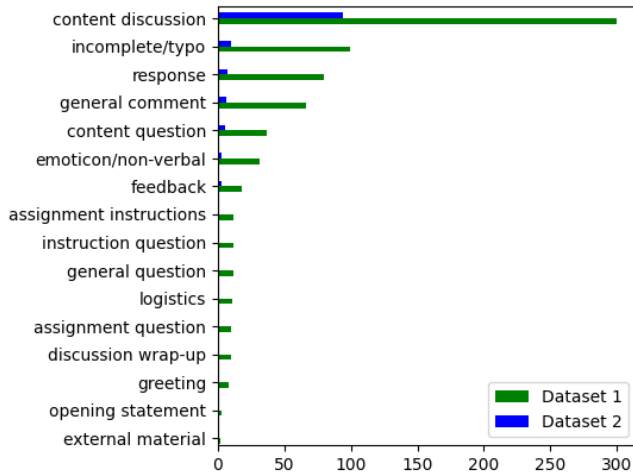


Figure 2. The distribution of message types in the dataset.




Figure 3. The distribution of active users in the dataset.

From all the graphs above we can see that in dataset 2, where there are messages concerning the discussion, we get a more uniform distribution of users as well as books. While the distribution of message types is dominated by the content discussion.

To better understand the contents we are trying to classify, we checked the top unigrams of the book texts. We also calculated the TF-IDF of the books as well as the collaborative discussions (as a whole) and displayed the top 15 stemmed words for each (see figure 4). This gives us a quick insight into the most important words of each book.



Figure 4. The top 15 words based on TF-IDF for each set.

From this we can say that the texts are contextually different, aside from a few words such as 'one', the top frequent words are more or less unique to their books.

Existing solutions

There are many approaches to automatic NLP text classification, which fall into three types of systems :

- rule-based systems,
- machine learning-based systems,
- hybrid systems.

Rule-based systems classify text into organized groups by using a set of handcrafted rules. But these systems require deep knowledge of the domain they are trying to classify. On the other hand, machine learning-based systems can learn the different associations between pieces of text by using pre-labeled examples as training data. Some of the more popular text classification algorithms [2] include Naive Bayes, Support Vector Machines, and deep learning. While hybrid systems apply both handcrafted rules and machine learning approaches. The current state-of-the-art consists of approaches such as BERT [3] or ELMo [4].

Handcrafted features

We decided to create handcrafted features as well and compare them to the performance of other feature extractors such as tf-idf and countvec. Unlike the other two we will not preprocess the text beforehand, meaning we will keep it as it is.

For our features we simply processed each message and created a vector of numbers, where the numbers in order represent :

- number of sentences,
- number of words,
- number of non-white characters,
- number of unique words,

- number of digits,
- number of stop words,
- number of interrogative words,
- the maximum word length,
- the average word length,
- the number of exclamation marks '!',
- the number of question marks '?',
- the number of commas ',',
- the number of dots '.',
- the number of dashes '-',
- the number of closing brackets ')',
- the number of opening brackets '(',
- the number of colons ':',
- the number of capital letters,
- the number of non-ascii characters.

Results

We then tested the all the combinations of models and extractors, in order to get a more in-depth analysis of the results. Note that RF is Random Forest, NB is Naive Bayes, LR is Logistic Regression, KNN is k-Nearest Neighbor and MV is Majority Voting.

Table 1. *CodePreliminary* with *tfidf* on dataset 1.

Model	Extractor	F1 Score	Accuracy
LR	countvec	0.52 ± 0.07	0.58 ± 0.06
NB	countvec	0.45 ± 0.08	0.52 ± 0.07
LR	handcrafted	0.44 ± 0.07	0.51 ± 0.06
RF	handcrafted	0.42 ± 0.07	0.51 ± 0.06
MV	handcrafted	0.42 ± 0.07	0.51 ± 0.06
LR	tfidf	0.39 ± 0.07	0.51 ± 0.06
RF	tfidf	0.39 ± 0.06	0.50 ± 0.05
MV	tfidf	0.39 ± 0.06	0.50 ± 0.05
KNN	handcrafted	0.39 ± 0.07	0.48 ± 0.07
NB	tfidf	0.34 ± 0.07	0.47 ± 0.06
RF	countvec	0.34 ± 0.07	0.47 ± 0.06
MV	countvec	0.34 ± 0.07	0.47 ± 0.06
NB	handcrafted	0.42 ± 0.06	0.46 ± 0.05
BERT	none	0.41 ± 0.21	0.42 ± 0.21
KNN	countvec	0.15 ± 0.07	0.13 ± 0.05
KNN	tfidf	0.07 ± 0.03	0.07 ± 0.03

Table 2. *Topic* with *tfidf* on dataset 1.

Model	Extractor	F1 Score	Accuracy
BERT	none	0.93 ± 0.07	0.92 ± 0.08
LR	countvec	0.77 ± 0.07	0.79 ± 0.05
NB	countvec	0.76 ± 0.06	0.77 ± 0.06
NB	tfidf	0.70 ± 0.06	0.75 ± 0.05
LR	tfidf	0.70 ± 0.05	0.74 ± 0.04
RF	tfidf	0.66 ± 0.07	0.71 ± 0.06
MV	tfidf	0.66 ± 0.07	0.71 ± 0.06
RF	countvec	0.64 ± 0.06	0.71 ± 0.05
MV	countvec	0.64 ± 0.06	0.71 ± 0.05
RF	handcrafted	0.57 ± 0.04	0.63 ± 0.04
MV	handcrafted	0.57 ± 0.04	0.63 ± 0.04
LR	handcrafted	0.55 ± 0.06	0.61 ± 0.06
KNN	countvec	0.50 ± 0.06	0.61 ± 0.05
KNN	tfidf	0.45 ± 0.05	0.59 ± 0.04
NB	handcrafted	0.56 ± 0.05	0.58 ± 0.04
KNN	handcrafted	0.50 ± 0.04	0.56 ± 0.04

Table 3. *BookID* with *tfidf* on dataset 1.

Model	Extractor	F1 Score	Accuracy
BERT	none	0.92 ± 0.07	0.92 ± 0.07
LR	countvec	0.77 ± 0.07	0.79 ± 0.05
NB	countvec	0.76 ± 0.06	0.77 ± 0.06
NB	tfidf	0.70 ± 0.06	0.75 ± 0.05
LR	tfidf	0.70 ± 0.05	0.74 ± 0.04
RF	tfidf	0.66 ± 0.07	0.71 ± 0.06
MV	tfidf	0.66 ± 0.07	0.71 ± 0.06
RF	countvec	0.64 ± 0.06	0.71 ± 0.05
MV	countvec	0.64 ± 0.06	0.71 ± 0.05
RF	handcrafted	0.57 ± 0.04	0.63 ± 0.04
MV	handcrafted	0.57 ± 0.04	0.63 ± 0.04
KNN	countvec	0.51 ± 0.06	0.62 ± 0.05
LR	handcrafted	0.55 ± 0.06	0.61 ± 0.05
KNN	tfidf	0.46 ± 0.04	0.59 ± 0.04
NB	handcrafted	0.56 ± 0.05	0.58 ± 0.04
KNN	handcrafted	0.50 ± 0.05	0.54 ± 0.05

Now on the second dataset, which contains only comments related to the books.

Table 4. *CodePreliminary* with *tfidf* on dataset 2.

Model	Extractor	F1 Score	Accuracy
NB	handcrafted	0.81 ± 0.11	0.82 ± 0.12
LR	countvec	0.79 ± 0.08	0.82 ± 0.10
LR	handcrafted	0.82 ± 0.09	0.82 ± 0.09
RF	handcrafted	0.79 ± 0.08	0.82 ± 0.07
MV	handcrafted	0.79 ± 0.08	0.82 ± 0.07
KNN	handcrafted	0.76 ± 0.10	0.80 ± 0.10
KNN	tfidf	0.67 ± 0.10	0.77 ± 0.08
RF	tfidf	0.61 ± 0.11	0.72 ± 0.09
NB	tfidf	0.61 ± 0.11	0.72 ± 0.09
MV	tfidf	0.61 ± 0.11	0.72 ± 0.09
LR	tfidf	0.61 ± 0.11	0.72 ± 0.09
RF	countvec	0.61 ± 0.11	0.72 ± 0.09
MV	countvec	0.61 ± 0.11	0.72 ± 0.09
NB	countvec	0.58 ± 0.09	0.62 ± 0.10
KNN	countvec	0.34 ± 0.15	0.29 ± 0.11
BERT	none	0.07 ± 0.17	0.12 ± 0.17

Table 5. *Topic* with *tfidf* on dataset 2.

Model	Extractor	F1 Score	Accuracy
LR	countvec	0.82 ± 0.11	0.82 ± 0.10
NB	countvec	0.81 ± 0.11	0.82 ± 0.11
LR	tfidf	0.80 ± 0.15	0.81 ± 0.13
BERT	none	0.79 ± 0.09	0.79 ± 0.07
NB	tfidf	0.77 ± 0.14	0.78 ± 0.13
RF	tfidf	0.71 ± 0.14	0.72 ± 0.14
MV	tfidf	0.71 ± 0.14	0.72 ± 0.14
RF	countvec	0.71 ± 0.14	0.72 ± 0.14
MV	countvec	0.71 ± 0.14	0.72 ± 0.14
KNN	tfidf	0.69 ± 0.18	0.70 ± 0.17
LR	handcrafted	0.55 ± 0.15	0.58 ± 0.13
RF	handcrafted	0.55 ± 0.14	0.57 ± 0.12
MV	handcrafted	0.55 ± 0.14	0.57 ± 0.12
NB	handcrafted	0.52 ± 0.15	0.52 ± 0.13
KNN	handcrafted	0.45 ± 0.16	0.46 ± 0.18
KNN	countvec	0.39 ± 0.17	0.45 ± 0.17

Table 6. *BookID* with *tfidf* on dataset 2.

Model	Extractor	F1 Score	Accuracy
NB	countvec	0.81 ± 0.11	0.82 ± 0.11
LR	countvec	0.82 ± 0.11	0.82 ± 0.10
BERT	none	0.81 ± 0.11	0.82 ± 0.10
LR	tfidf	0.80 ± 0.15	0.81 ± 0.13
NB	tfidf	0.77 ± 0.14	0.78 ± 0.13
RF	tfidf	0.71 ± 0.14	0.72 ± 0.14
MV	tfidf	0.71 ± 0.14	0.72 ± 0.14
RF	countvec	0.71 ± 0.14	0.72 ± 0.14
MV	countvec	0.71 ± 0.14	0.72 ± 0.14
KNN	tfidf	0.67 ± 0.18	0.68 ± 0.18
LR	handcrafted	0.55 ± 0.15	0.58 ± 0.13
RF	handcrafted	0.55 ± 0.14	0.57 ± 0.12
MV	handcrafted	0.55 ± 0.14	0.57 ± 0.12
NB	handcrafted	0.52 ± 0.15	0.52 ± 0.13
KNN	handcrafted	0.44 ± 0.15	0.44 ± 0.17
KNN	countvec	0.37 ± 0.16	0.44 ± 0.15

Conclusion

In our assignment we compared multiple classifiers and models in IMapBooks discussion classification.

When comparing the different extractors we came to a conclusion that the countvec performed the best, and then depending on the classification target handcrafted and tf-idf. When we classified CodePreliminary the handcrafted features performed better than the tf-idf, but performed worse when we classified Book ID and Topic. This is due to the fact that our handcrafted features only take account the structure of the messages, which would make more sense for CodePreliminary.

Overall BERT performed great with an accuracy of 92 % on the first dataset. An interesting observation is that for some reason BERT performed very poorly when predicting the CodePreliminary. Unfortunately we ran out of time to do a more in-depth analysis there. It is important to also note that BERT's accuracy dropped in the second dataset, while the accuracy of other classifiers increased. This is due to a small dataset, which is not appropriate to use for learning deeper models. It still performed decent on Book ID and Topic, but due to the time complexity it took to predict, simpler classifiers, which take almost no time are more appropriate here.

The results also showed that most of the time Logistic Regression performed best out of the classifiers. An important result was also that KNN performed significantly better on the dataset containing only messages pertaining to the discussion.

All this means that proper filtering is important when using text classification, as any unwanted or not relevant

data will result in significant losses to the accuracy. Also if we use deeper networks, we definitely need enough data to train it, otherwise the results will not be good, due to overfitting or not learning at all.

References

- [1] Grandon Gill and Glenn Smith. Imapbook: Engaging young readers with games. *Journal of Information Technology Education: Discussion Cases*, 2:10, 01 2013.
- [2] Tomas Prankevicius and V. Marcinkevicius. Comparison of naive bayes, random forest, decision tree, support vector machines, and logistic regression classifiers for text reviews classification. *Balt. J. Mod. Comput.*, 5, 2017.
- [3] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. 2019.
- [4] Matthew E. Peters, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, and Luke Zettlemoyer. Deep contextualized word representations, 2018.